

Kaggle Heart Attack Data Analysis Prediction/Classification using Logistic Regression, Random Forest

Rohit Gupta

July 26, 2021

```
In [1]: # By : Rohit Gupta, Date : Date : 12 July 2021
# Reading and Analyzing covid data from the following kaggle dataset :
# https://www.kaggle.com/rashikrahmanpritom/heart-attack-analysis-prediction-dataset

import os

In [32]: os.getcwd()
os.chdir("C:\\\\Users\\\\....\\\\Desktop\\\\Self_Learning\\\\heart_dataset")

In [33]: import os
path = "C:\\\\Users\\\\....\\\\Desktop\\\\Self_Learning\\\\heart_dataset"
arr = os.listdir(path)
print(arr)

['heart.csv', 'o2Saturation.csv']

In [34]: import pandas as pd
data_heart = pd.read_csv('heart.csv')
data_heart.head()

Out[34]:    age  sex  cp  trtbps  chol  fbs  restecg  thalachh  exng  oldpeak  slp \
0    63    1    3    145   233    1      0     150      0      2.3    0
1    37    1    2    130   250    0      1     187      0      3.5    0
2    41    0    1    130   204    0      0     172      0      1.4    2
3    56    1    1    120   236    0      1     178      0      0.8    2
4    57    0    0    120   354    0      1     163      1      0.6    2

      caa  thall  output
0      0      1      1
1      0      2      1
2      0      2      1
3      0      2      1
4      0      2      1
```

```
In [5]: data_heart.columns
```

```
Out[5]: Index([u'age', u'sex', u'cp', u'trtbps', u'chol', u'fbs', u'restecg',  
    u'thalachh', u'exng', u'oldpeak', u'slp', u'caa', u'thall', u'output'],  
    dtype='object')
```

```
In [6]: # What does the 1 and 0 stand in the sex column ? Assume 1 stands for male and 0 stands  
  
data_heart.shape
```

```
Out[6]: (303, 14)
```

```
In [7]: # what kind of exploratory data analysis can we perform here ? Let us see  
# the correlation between Age, Sex, exang: exercise induced angina and trtbps : resting  
# Possibility that with age, and increased stress levels of life, resting blood pressure  
import seaborn as sns  
import matplotlib.pyplot as plt  
  
# allow plots to occur inline  
%matplotlib inline
```

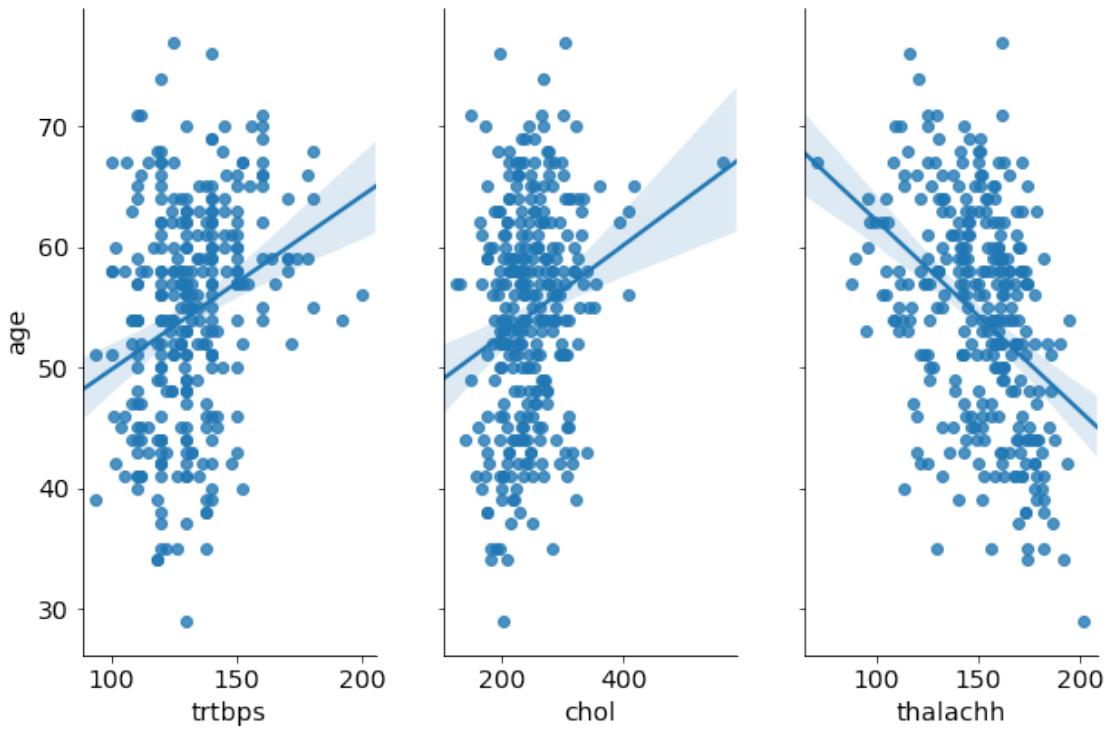
```
In [8]: data_heart.sort_values(by=[ 'age'],axis=0,ascending=True).head()
```

```
Out[8]:      age  sex  cp  trtbps  chol  fbs  restecg  thalachh  exng  oldpeak  slp  \  
72      29     1     1      130    204     0        0      202     0      0.0      2  
58      34     1     3      118    182     0        0      174     0      0.0      2  
125     34     0     1      118    210     0        1      192     0      0.7      2  
239     35     1     0      126    282     0        0      156     1      0.0      2  
65      35     0     0      138    183     0        1      182     0      1.4      2  
  
      caa  thall  output  
72      0     2      1  
58      0     2      1  
125     0     2      1  
239     0     3      0  
65      0     2      1
```

```
In [65]: # Using just scatter plots we do not get much insight. However we can use kind = 'reg'  
# are somewhat correlated.  
# sns.pairplot(data_heart, x_vars=['trtbps', 'chol', 'cp'], y_vars=['age'], aspect = 0.5)
```

```
In [228]: # What other plots can be drawn that provide other descriptive insights into the data?  
sns.pairplot(data_heart, x_vars=['trtbps', 'chol', 'thalachh'], y_vars=['age'], aspect=1)
```

```
Out[228]: <seaborn.axisgrid.PairGrid at 0x16738518>
```

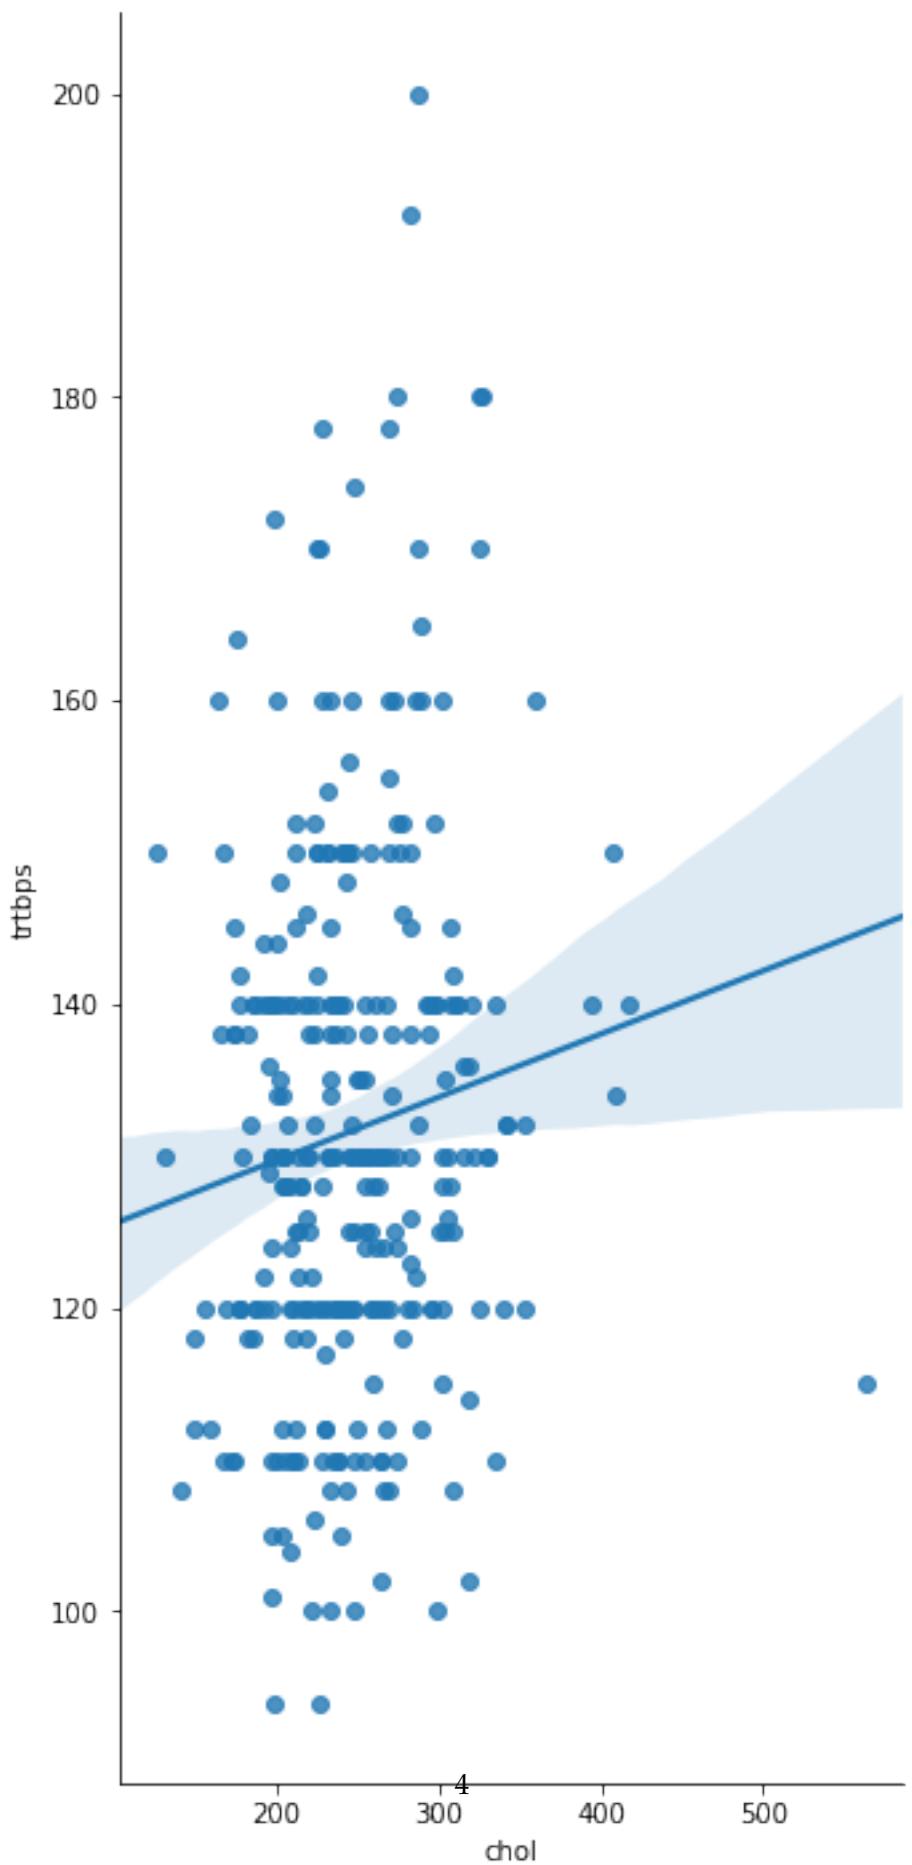


```
In [64]: # It will be interesting to see how the variables age and rest_ecg are related.
# sns.pairplot(data_heart, x_vars = ['restecg'], y_vars = ['age'], height = 6 ,aspect =
```

```
In [24]: # Further see how other variables are related to each other
# how is resting blood pressure associated with cholesterol levels ?
# In the process determine which features will predict heart attack.
```

```
sns.pairplot(data_heart, x_vars = ['chol'], y_vars = ['trtbp'], height = 10 ,aspect =
```

```
Out[24]: <seaborn.axisgrid.PairGrid at 0xe987898>
```



```
In [25]: # Draw boxplots to see the measures of the features - since there are several continuous
# using boxplot we can get the distribution of the data.
#help(sns.boxplot)

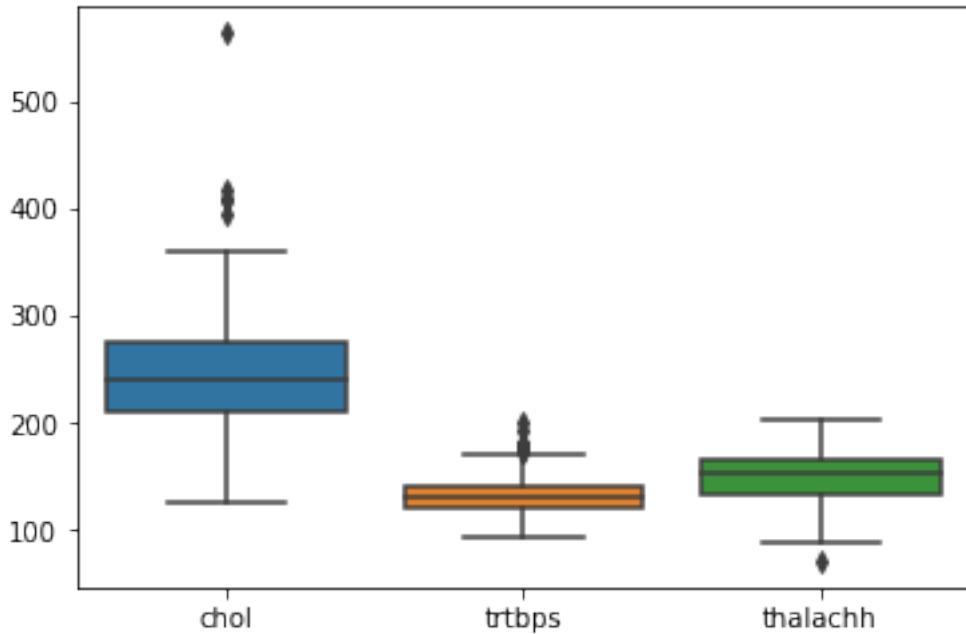
feature_cols = ['chol','trtbps','thalachh']
sns.boxplot(data=data_heart[feature_cols])

## Grouped boxplots for Male and Female :

# select all rows with Sex = 0 ---> Male ----> df1
# select all rows with Sex = 1 ---> Female ---> df 2

# then compare boxplots across the chol, trtbps and thalachh
```

Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0xf1c6d30>



```
In [26]: data_heart_male = data_heart.loc[data_heart['sex']==0]
data_heart_female = data_heart.loc[data_heart['sex']==1]
```

In [27]: data_heart_male.shape

Out[27]: (96, 14)

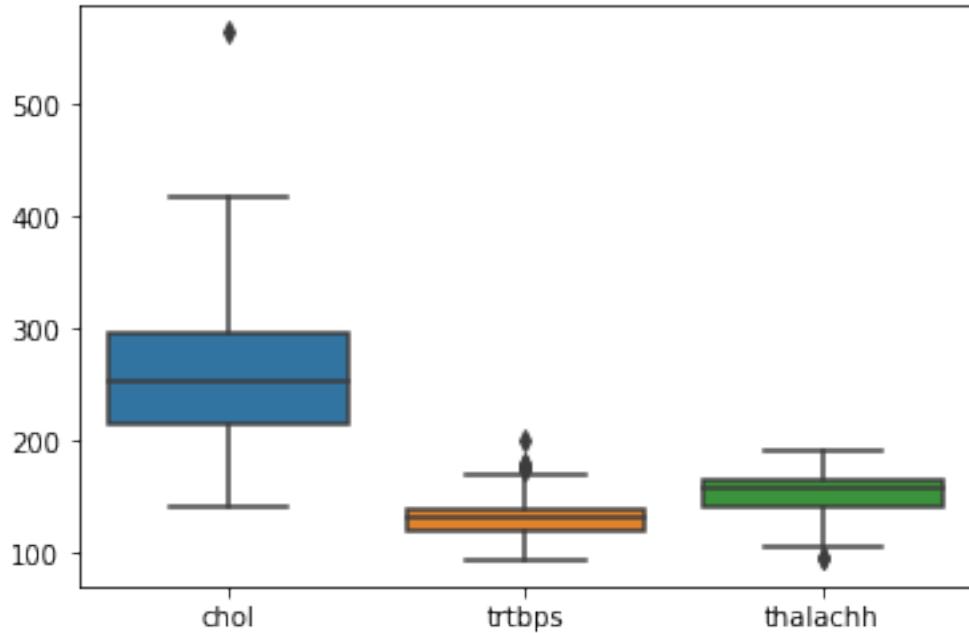
In [28]: data_heart_female.shape

Out[28]: (207, 14)

In [17]: # Category Male boxplot

```
feature_cols = ['chol','trtbps','thalachh']
sns.boxplot(data=data_heart_male[feature_cols])
```

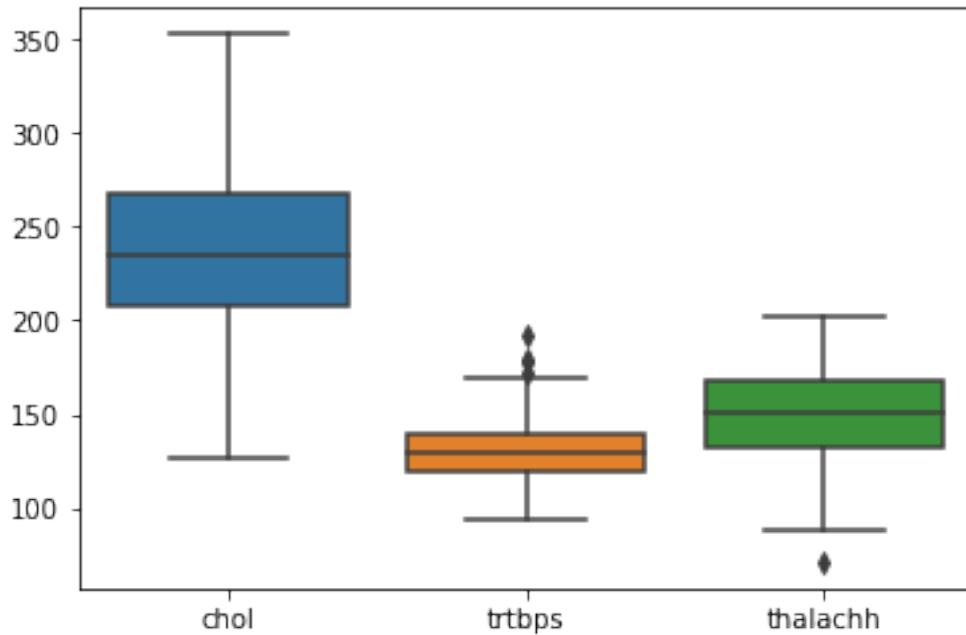
Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0xe1337f0>



In [18]: # Category Female boxplots

```
feature_cols = ['chol','trtbps','thalachh']
sns.boxplot(data=data_heart_female[feature_cols])
```

Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0xe3a3470>

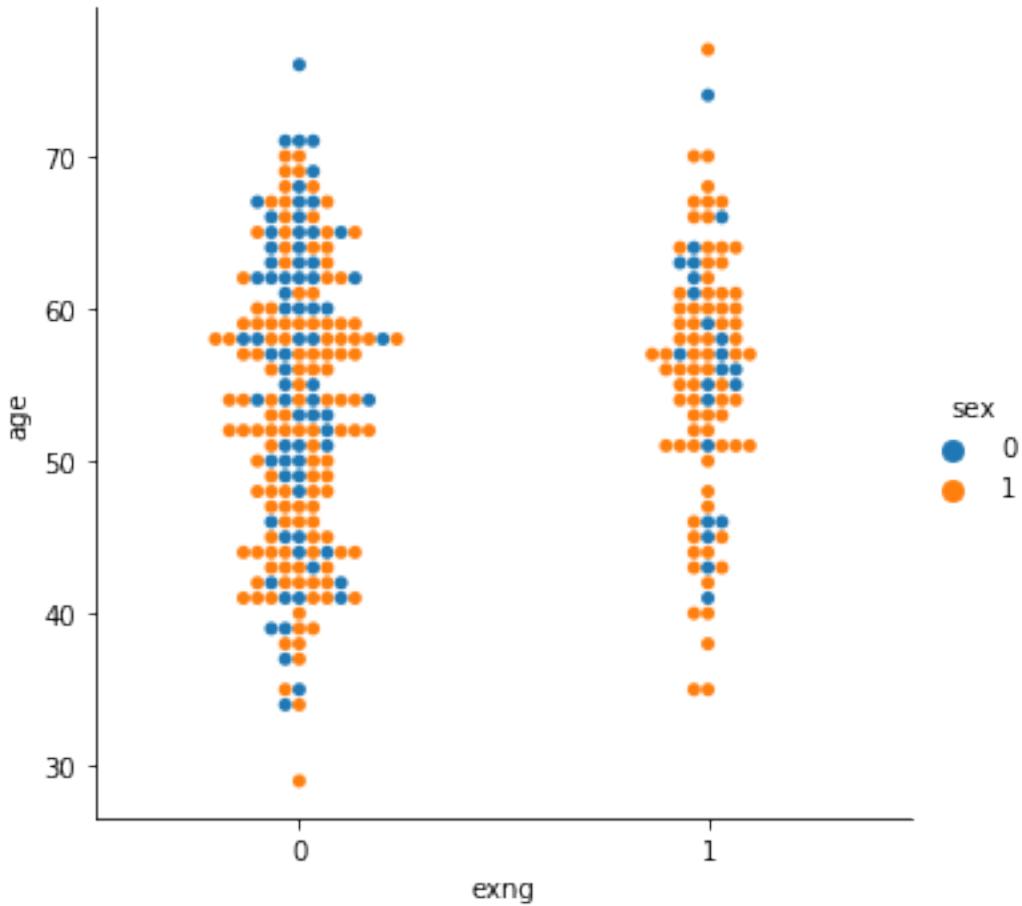


```
In [43]: # Visualizing the given categorical data such as :
# exang: exercise induced angina (1 = yes; 0 = no)
# Recall 0 : stands for male and 1 stands for female

sns.catplot(x="exng", y="age", hue="sex", kind="swarm", data=data_heart)

# We see that excercise induced angina more likely in females than in males.
# Excercise induced angina less in males (as compared to females)
```

Out[43]: <seaborn.axisgrid.FacetGrid at 0x11733240>



In [59]: # Visualizing cp vs age : with males/females :

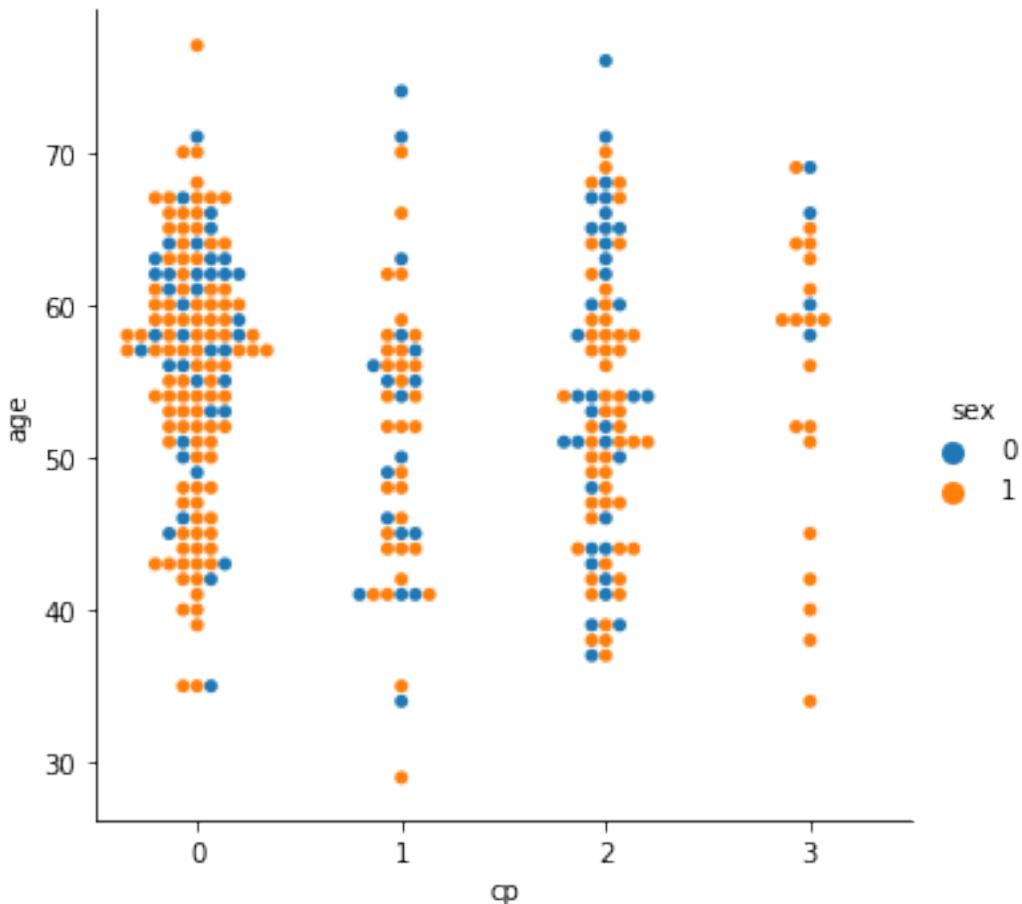
```
# cp : Chest Pain type chest pain type

# Value 1: typical angina
# Value 2: atypical angina
# Value 3: non-anginal pain
# Value 4: asymptomatic

sns.catplot(x="cp", y="age", hue="sex", kind="swarm", data=data_heart)

# vertical positioning of the categorical labels :
#sns.catplot(x="age", y="cp", hue="sex", kind="swarm", data=data_heart)
```

Out[59]: <seaborn.axisgrid.FacetGrid at 0x13eff240>

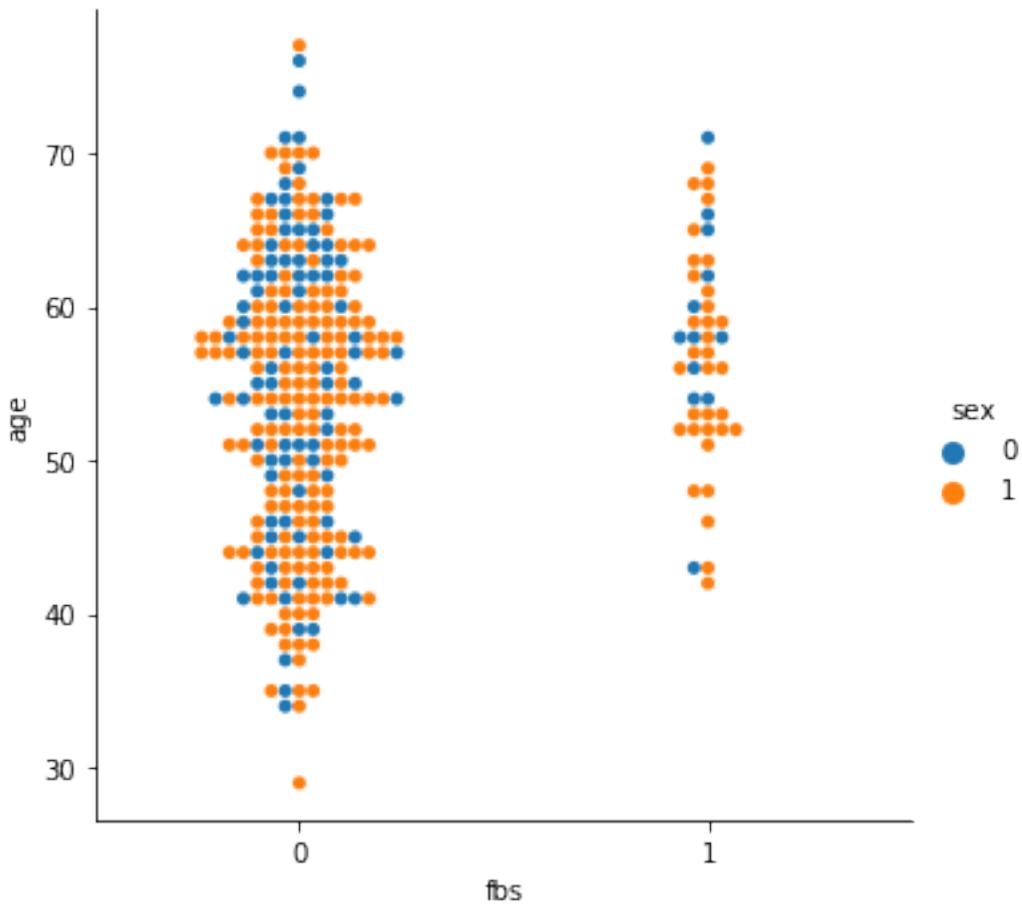


```
In [56]: # Visualizing fasting blood sugar in males vs females with age on x-axis :
# fbs : (fasting blood sugar > 120 mg/dl) (1 = true; 0 = false)

sns.catplot(x="fbs", y="age", hue="sex", kind = "swarm", data=data_heart)

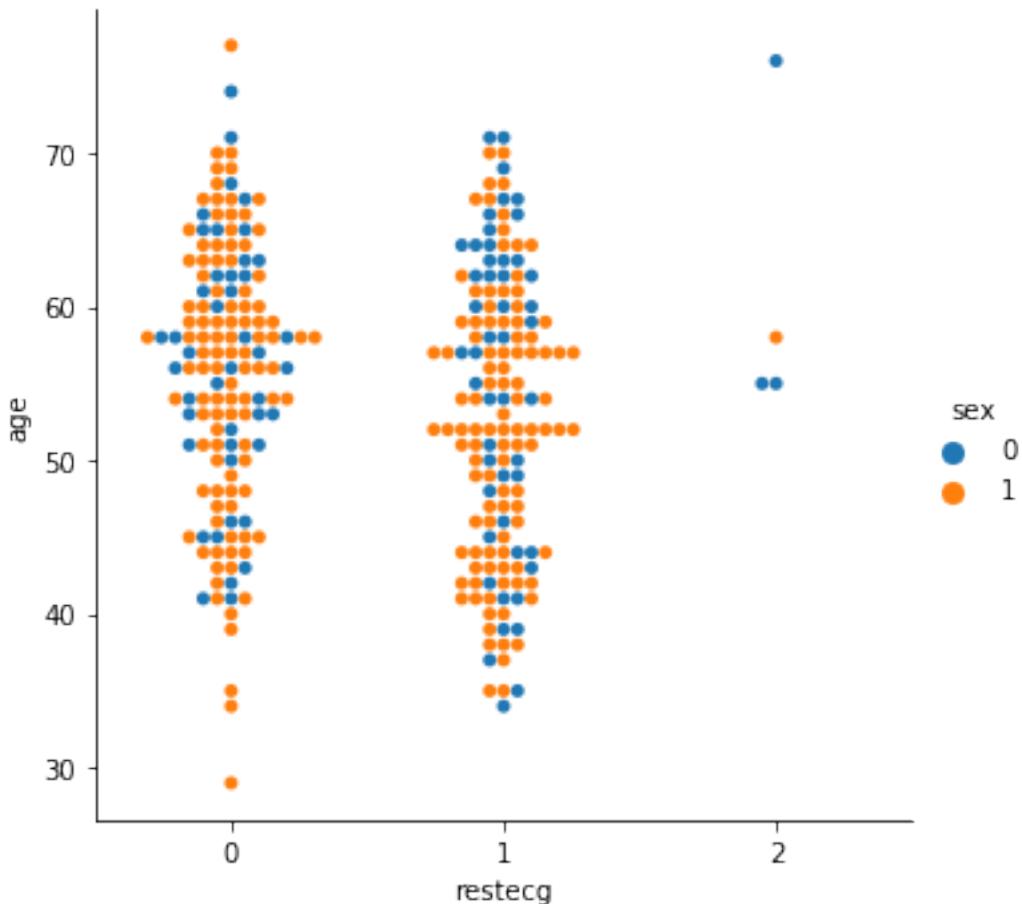
# majority of the population resides in the fasting blood sugar < 120 mg/dL
```

Out[56]: <seaborn.axisgrid.FacetGrid at 0x12e657b8>



```
In [61]: # visualizing the rest_ecg with respect to age  
  
# rest_ecg : resting electrocardiographic results  
  
# Value 0: normal  
# Value 1: having ST-T wave abnormality (T wave inversions and/or ST elevation or depression)  
# Value 2: showing probable or definite left ventricular hypertrophy by Estes' criteria  
  
sns.catplot(x="restecg" , y = "age" , hue = "sex" , kind = "swarm" , data=data_heart)
```

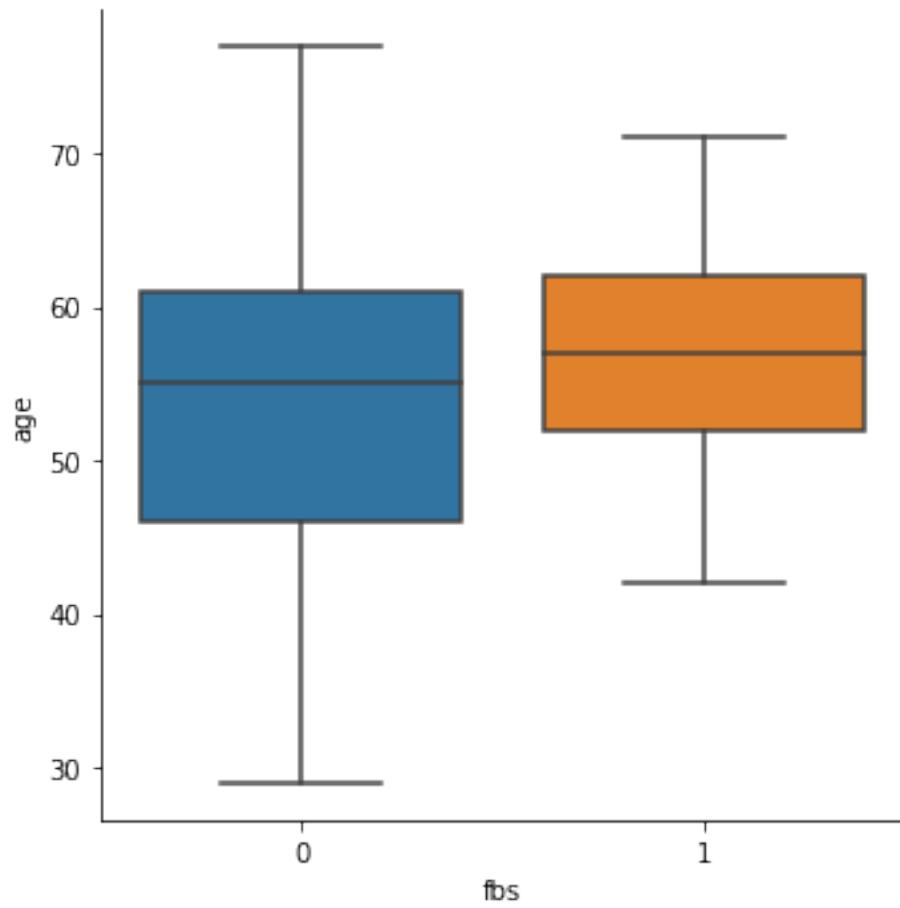
Out[61]: <seaborn.axisgrid.FacetGrid at 0x140a8240>



```
In [67]: # fbs boxplot
sns.catplot(x="fbs", y="age", kind="box", data=data_heart)

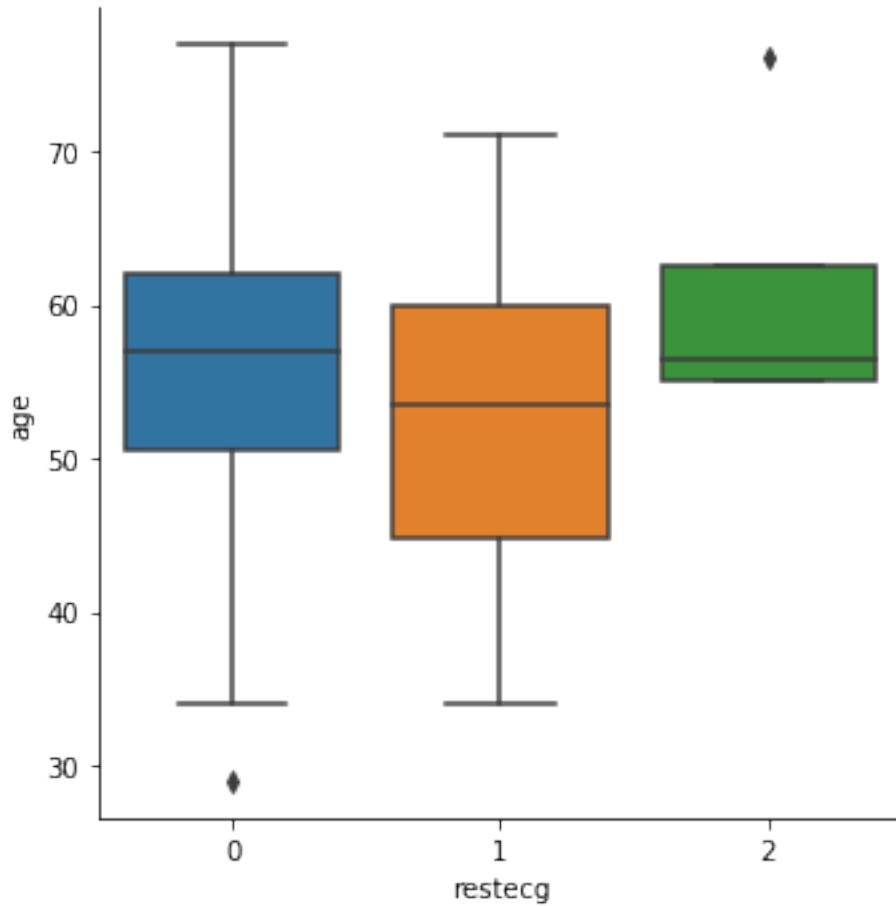
# fbs with respect to males
# sns.catplot(x="fbs", y="age", kind="box", data=data_heart_male)
# sns.catplot(x="fbs", y="age", kind="box", data=data_heart_female)
```

Out[67]: <seaborn.axisgrid.FacetGrid at 0x132eef60>



In [74]: `sns.catplot(x="restecg", y="age", kind="box", data=data_heart)`

Out[74]: <seaborn.axisgrid.FacetGrid at 0x152d6be0>



```
In [66]: # correlation plot between the features/variables :
```

```
Out[66]: Index([u'age', u'sex', u'cp', u'trbps', u'chol', u'fbs', u'restecg',
       u'thalachh', u'exng', u'oldpeak', u'slp', u'caa', u'thall', u'output'],
      dtype='object')
```

```
In [241]: # Predicting heart risk based on these features :
```

```
# cholesterol, resting blood pressure and maximum heart rate achieved
```

```
#data_heart.columns
#data_heart.head(40)
```

```
In [75]: # Prediction/Classification of heart attack on the basis of following features :
```

```
feature_cols = ['trbps', 'chol', 'thalachh']
target_cols = ['output']
```

```
X = data_heart[feature_cols]
# recall y has to be a pandas series
```

```

Y = data_heart.output

from sklearn.cross_validation import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, random_state=1)

# import model
from sklearn.linear_model import LogisticRegression

# instantiate model
logreg = LogisticRegression()

C:\ProgramData\Anaconda2\lib\site-packages\sklearn\cross_validation.py:41: DeprecationWarning: T
  "This module will be removed in 0.20.", DeprecationWarning)

In [76]: #X_train.shape
          #Y_train.shape
          #X_test.shape
          #Y_test.shape

          # fit the model on the data
logreg.fit(X_train, Y_train)

Out[76]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)

In [77]: # perform prediction
Y_pred = logreg.predict(X_test)

In [78]: # calculate RMSE error between Y_test and Y_pred
          # and see how it changes if you remove one feature from the feature_cols

          from sklearn import metrics
          import numpy as np
          # print(np.sqrt(metrics.mean_squared_error(Y_test,Y_pred)))
          # 0.5619514869490163

          print(metrics.accuracy_score(Y_test,Y_pred))

0.6842105263157895

In [79]: # remove chol and see how the rmse value changes :
feature_cols = ['trtbps', 'thalachh', 'age']

X = data_heart[feature_cols]
Y = data_heart.output

```

```

# train-test split
from sklearn.cross_validation import train_test_split

X_train, X_test, Y_train, Y_test = train_test_split(X,Y,random_state=1)

# import
from sklearn.linear_model import LogisticRegression

# instantiate
logreg = LogisticRegression()

In [80]: #fit the model
logreg.fit(X_train,Y_train)

Out[80]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, max_iter=100, multi_class='ovr', n_jobs=1,
                             penalty='l2', random_state=None, solver='liblinear', tol=0.0001,
                             verbose=0, warm_start=False)

In [81]: Y_pred_new = logreg.predict(X_test)

In [165]: # For the logistic regression model print the confusion matrix
metrics.confusion_matrix(Y_test,Y_pred)

Out[165]: array([[21, 14],
                 [10, 31]], dtype=int64)

In [223]: # splice the confusion matrix into 4 parts : TP,TN,FP,FN

confusion_heartmat = metrics.confusion_matrix(Y_test,Y_pred)
TP = confusion_heartmat[1,1]
TN = confusion_heartmat[0,0]
FP = confusion_heartmat[0,1]
FN = confusion_heartmat[1,0]

In [168]: # calculate RMSE error between Y_test and Y_pred_new
# print(np.sqrt(metrics.mean_squared_error(Y_test,Y_pred_new)))
# 0.5735393346764044

print(metrics.accuracy_score(Y_test,Y_pred_new))

0.6710526315789473

In [227]: # accuracy : TP+TN/TP+TN+FP+FN
print(1-metrics.accuracy_score(Y_test,Y_pred))

0.3157894736842105

```

```
In [220]: # when actual value is positive, how often is the prediction correct :  
#TP/TP+FN
```

```
# TPR = sensitivity = Recall :  
metrics.recall_score(Y_test,Y_pred)
```

```
Out[220]: 0.7560975609756098
```

```
In [224]: # FPR - specificity  
print(FP/float(FP+TN))
```

```
0.4
```

```
In [226]: from sklearn.metrics import classification_report  
print(classification_report(Y_test, Y_pred, labels=[0, 1]))
```

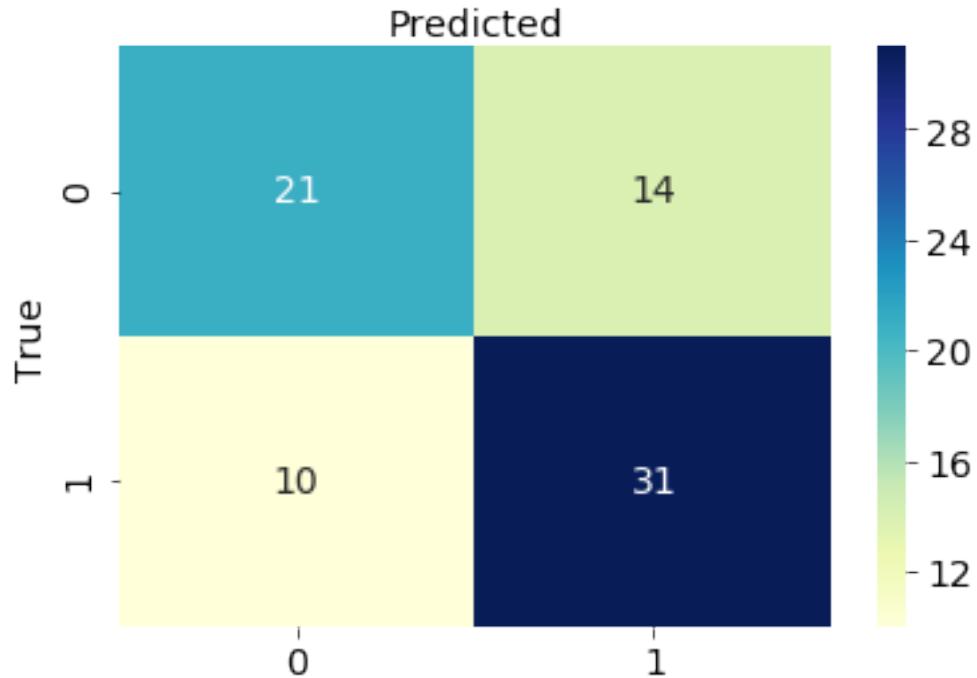
	precision	recall	f1-score	support
0	0.68	0.60	0.64	35
1	0.69	0.76	0.72	41
avg / total	0.68	0.68	0.68	76

```
In [242]: conf_matrix_logreg_model = metrics.confusion_matrix(Y_test, Y_pred)
```

```
labels = [0, 1]  
fig, ax = plt.subplots()  
tick_marks = np.arange(len(labels))  
plt.xticks(tick_marks, labels)  
plt.yticks(tick_marks, labels)  
  
# create heatmap  
sns.heatmap(pd.DataFrame(conf_matrix_logreg_model), annot=True, cmap="YlGnBu", fmt='g')  
ax.xaxis.set_label_position("top")  
plt.title('Confusion Matrix for Logistic Regression Model', y=1.1)  
plt.ylabel('True')  
plt.xlabel('Predicted')
```

```
Out[242]: Text(0.5,11,'Predicted')
```

Confusion Matrix for Logistic Regression Model



```
In [216]: # Draw the ROC-AUC-Curve :
```

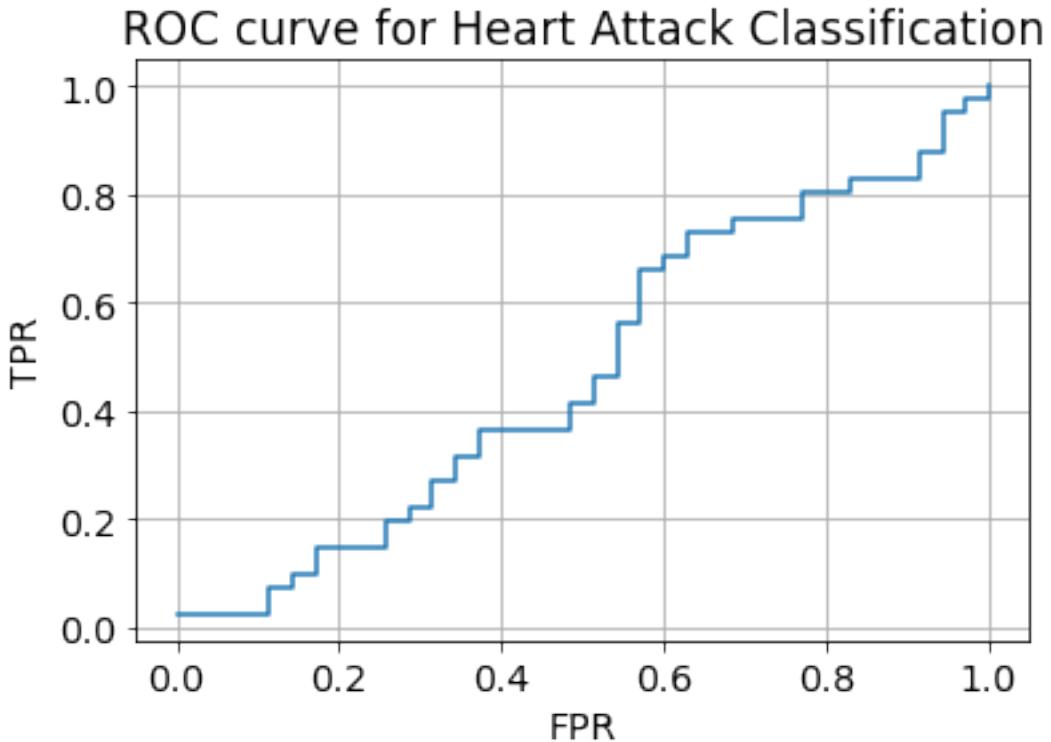
```
# calculate Y_pred_prob
Y_pred_prob = logreg.predict_proba(X_test)[:,1]
#Y_test.shape
#Y_pred_prob.shape

fpr, tpr, threshold = metrics.roc_curve(Y_test,Y_pred_prob)
```

```
In [219]: # allow plots to appear in the notebook
```

```
%matplotlib inline
import matplotlib.pyplot as plt
plt.rcParams['font.size'] = 14

plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC curve for Heart Attack Classification')
plt.grid(True)
```



```
In [232]: # Let us use a different model to see if the accuracy can be improved :
# Let us use random forest classifier :
```

```
from sklearn.ensemble import RandomForestClassifier

# instantiate
rf_model = RandomForestClassifier(max_depth=2, random_state=0)

feature_cols = ['trtbps', 'chol', 'thalachh']
#target_cols = ['output']

X = data_heart[feature_cols]
# recall y has to be a pandas series
Y = data_heart.output

from sklearn.cross_validation import train_test_split
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, random_state=1)
```

```
In [233]: # fit the model
rf_model.fit(X_train, Y_train)

# store prediction values
Y_pred_rf_model = rf_model.predict(X_test)
```

```
# test the accuracy score
print(metrics.accuracy_score(Y_test,Y_pred_rf_model))
```

```
0.6973684210526315
```

In [248]: # use RandomizedSearchCV with random forest to tune the model to perform better :
print(rf_model.get_params)

```
from sklearn.grid_search import RandomizedSearchCV
# specify "parameter distributions" rather than a "parameter grid"

max_values = range(2,40)
estimators = range(2,20)

param_dist = dict(max_depth = max_values, n_estimators=estimators)

# n_iter controls the number of searches
# instantiate the RandomizedSearchCV model
rf_randcv_model = RandomizedSearchCV(rf_model,param_dist,cv=10,scoring='accuracy',n_it

# train the model
rf_randcv_model.fit(X_train,Y_train)

<bound method RandomForestClassifier.get_params of RandomForestClassifier(bootstrap=True, class_
    max_depth=2, max_features='auto', max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
    oob_score=False, random_state=0, verbose=0, warm_start=False)>
```

Out[248]: <bound method RandomizedSearchCV.get_params of RandomizedSearchCV(cv=10, error_score=''
 estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterio
 max_depth=2, max_features='auto', max_leaf_nodes=None,
 min_impurity_decrease=0.0, min_impurity_split=None,
 min_samples_leaf=1, min_samples_split=2,
 min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
 oob_score=False, random_state=0, verbose=0, warm_start=False),
 fit_params={}, iid=True, n_iter=10, n_jobs=1,
 param_distributions={'n_estimators': [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13
 pre_dispatch='2*n_jobs', random_state=5, refit=True,
 scoring='accuracy', verbose=0)>

In [251]: # get best parameters
rf_randcv_model.best_params_

get best scores:
rf_randcv_model.best_score_

```
Out[251]: 0.6696035242290749
```

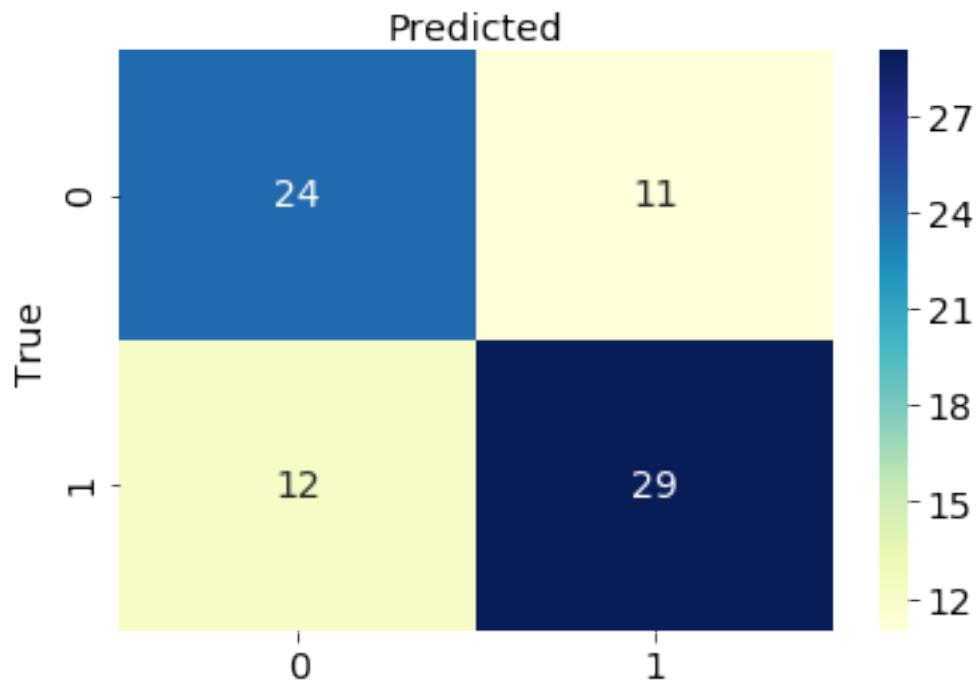
```
In [237]: conf_matrix_rf_model = metrics.confusion_matrix(Y_test, Y_pred_rf_model)

labels = [0, 1]
fig, ax = plt.subplots()
tick_marks = np.arange(len(labels))
plt.xticks(tick_marks, labels)
plt.yticks(tick_marks, labels)

# create heatmap
sns.heatmap(pd.DataFrame(conf_matrix_rf_model), annot=True, cmap="YlGnBu", fmt='g')
ax.xaxis.set_label_position("top")
plt.title('Confusion Matrix for Random Forest Model', y=1.1)
plt.ylabel('True')
plt.xlabel('Predicted')

Out[237]: Text(0.5,11,'Predicted')
```

Confusion Matrix for Random Forest Model



```
In [235]: metrics.confusion_matrix(Y_test, Y_pred_rf_model)
```

```
Out[235]: array([[24, 11],
 [12, 29]], dtype=int64)
```

```
In [246]: # # Implement NaiveBayes and check if accuracy can be improved or not.
# from sklearn.naive_bayes import GaussianNB

# # instantiate model
# nb_model = GaussianNB()

# # perform training
# nb_model.fit(X_train, Y_train)

# # perform testing
# Y_pred_nb_model = nb_model.predict(X_test)

# # check accuracy
# print(metrics.accuracy_score(Y_test, Y_pred_nb_model))

# #0.631578947368421
```

0.631578947368421

In [163]: # Future Work - Additional Improvements for accuracy to be tested in future :

```
# # Use linear regression with cross validation to build model, do prediction and test

# from sklearn.linear_model import LinearRegression
# linreg = LinearRegression()

# from sklearn.cross_validation import cross_val_score
# mse_scores = cross_val_score(linreg, X, Y, cv=10, scoring = 'mean_squared_error')

# pos_mse_scores = -mse_scores
# # print(pos_mse_scores)
# import numpy as np
# rmse_scores = np.sqrt(pos_mse_scores)
# print(rmse_scores.mean())

# # test model by removing cholesterol levels
# feature_cols = ['trtbps', 'thalachh']
# #target_cols = ['output']

# X = data_heart[feature_cols]
# # recall y has to be a pandas series
# Y = data_heart.output

# mse_scores_2 = cross_val_score(linreg, X, Y, cv=10, scoring = 'mean_squared_error')
# print(mse_scores_2)

# pos_mse_scores_2 = -mse_scores_2
```

```

# rmse_scores_2 = np.sqrt(pos_mse_scores_2)
# print(rmse_scores_2.mean())

# # Use KNN for prediction

# from sklearn.neighbors import KNeighborsClassifier

# # find optimal value of k between 1 to 31 for this

# k_values = range(1,31)
# all_scores = []
# for k in k_values :
#     # instantiate
#     knn = KNeighborsClassifier(n_neighbors=k)
#     # score calculation
#     scores = cross_val_score(knn,X,Y,cv=10,scoring='accuracy')
#     # append
#     all_scores.append(scores.mean())

# # Very Very Very Intersting : I am achieving an accuracy of 70 % when input to cross
# import numpy as np
# #np.mean(all_scores)
# #print(all_scores)
# max(all_scores)

# print(all_scores)

# max(all_scores)

# # plot K-value corresponding to accuracy
# import matplotlib.pyplot as plt

# plt.plot(k_values, all_scores)
# plt.title('K_Values Vs all_scores')
# plt.xlabel('k_values')
# plt.ylabel('all_scores')
# plt.show()

# # n_neighbors = 29

# knn = KNeighborsClassifier(n_neighbors=29)
# knn.fit(X_train,Y_train)
# Y_pred_best_knn_model = knn.predict(X_test)
# print(metrics.accuracy_score(Y_test,Y_pred_best_knn_model))

# ## Accuracy decreased for some reason.

```

